

Literature Review of Optimisation Techniques to the Graph Colouring Problem

Scott Whitemore & Brendon Veronese

October 22, 2015

Abstract

A review of papers concerning the current state-of-the-art for solving the GCP is presented. Algorithms from these papers are implemented and compared, results are shown in an unbiased way.

Chapter 1

Literature Review

1.1 Introduction

1.2 Johnson - Simulated Annealing

1.3 Gravitational Swarm Intelligence

The natural inspiration of this algorithm does not come from living beings, such as ants or bees, but from the basic physical law of gravitational attraction between objects. We construct a world that agents navigate through, attracted by the gravitational pull of specific objects, the color goals, such that they may suffer specific repulsion forces, activated by the friend-or-foe nature of the relation between agents induced by the adjacency relation in the underlying graph.

Initial definition: Let $G = (V, E)$ be a graph defined on a set of nodes $V = \{v_1, \dots, v_N\}$ and edges $E \subseteq V \times V$. We define a group of GS-GC agents $B = \{b_1, b_2, \dots, b_N\}$ each corresponding to a graph node. Each agent navigates inside a square planar toric world according to a speed vector \vec{v}_i . At any moment in time we know the position attribute of each agent $p_i(t) = (x_i, y_i)$ where x_i and y_i are the Cartesian coordinates in the space. When $t = 0$ we have the initial position of the agents $p_i(0) = (x_0, y_0)$. Suppose that we want to color the graph with K colors, denoting $C = \{1, 2, \dots, K\}$ the set of colors, where K must not be lower than the chromatic number of the graph for the GS-GC to converge. We assign to these colors, K fixed points in space, the color goals $CG = \{g_1, \dots, g_K\}$, endowed with a gravitational attraction resulting in a velocity component \vec{v}_{gc} affecting the agents. The attraction force decreases with the distance, but affects all the agents in the space.

The problem collapses into the minimisation of a cost function:

$$\min |\{b_i \text{ s.t. } b_i \in \{g_1, \dots, g_N\}\}|$$

We denote the set of agents whose position is in the region of the space near enough to a color neighbourhood of the color as:

$$\mathcal{N}(g_k) = \{b_i \text{ s.t. } \|p_i - g_k\| < nearenough\} \quad (1.1)$$

We denote the fact that the node has been assigned to the corresponding color assigning value to a the agent color attribute

$$b_i \in \mathcal{N}(g_k) \Rightarrow c_i = k \quad (1.2)$$

The initial value of the agent color attribute c_i is zero or null. Inside the spatial neighbourhood of a color goal there is no further gravitational attraction. However, there may be a repulsion force between agents that are connected with an edge in the graph G . This repulsion is only effective for agents inside the same color goal neighbourhood. To model this effect, we define function repulsion which has value 1 if a pair of GS-GC agents have an edge between them, and 0 otherwise. The repulsive forces experimented by agent b_i from the agents in the color goal g_k are computed as follows:

$$R(b_i, g_k) = \sum_{b_j \in \mathcal{N}(g_k)} \text{repulsion}(b_i, b_j) \quad (1.3)$$

The cost function defined on the global system spatial configuration is:

$$f(B, CG) = |\{b_i \text{ s.t. } c_i \in C \& R(b_i, g_{c_i}) = 0\}| \quad (1.4)$$

This cost function is the number of graph nodes which have a color assigned and no conflict inside the color goal. The agents outside the neighbourhood of any color goal can't be evaluated, so it can be a part of the solution of the problem. The dimension of the world and the definition of the *nearenough* threshold allows controlling the speed of convergence of the algorithm. If the world is big and the *nearenough* variable is small then the algorithm converges slowly but monotonically to the solution, if the world is small and the *nearenough* variable is big the algorithm is faster but convergence is jumpy because the algorithm falls in local minima and needs transitory energy increases to escape them. The reason of this behaviour is that the world is not normalized and the magnitude of the velocity vector can be bigger than the color goal spatial influence and can cross a goal without falling in it.

Each color goal has an attraction well spanning the entire space, therefore the gravitational analogy. But in our approach the magnitude of the attraction drops proportionally with the Euclidean distance d between the goal and the GS-GC agent, but it never disappears. If $\|d\| < \text{nearenough}$ then we make $d = 0$, and the agent's velocity becomes 0 stopping it.

Definition: A Gravitational Swarm (GS) is a collection of particles $P = \{p_1, \dots, p_L\}$ moving in an space S subjected to attraction and repulsion forces. Attraction correspond to long range gravitational interactions. Repulsions correspond to short range electrical interactions. Particle attributes are: spatial localization $s_i \in S$, mass $m_i \in \mathbb{R}$, charge $\mu_i \in \mathbb{R}$, set of repelled particles $r_i \subseteq P$. The motion of the particle in the space is governed by equation:

$$s_i(t) = -m_i(t) A_i(t) + \mu_i(t) R_i(t) + \eta(t) \quad (1.5)$$

where $A_i(t)$ and $R_i(t)$ are the result of the attractive and repulsive forces, and $\eta(t)$ is a random (small) noise term. The attractive motion term is of the form:

$$A_i(t) = \sum_{p_j \in P - r_i} m_j(t)(s_i - s_j) \delta_{ij}^A \quad (1.6)$$

where

$$\delta_{ij}^A = \begin{cases} \|s_i - s_j\|^{-2} & \|s_i - s_j\|^2 > \theta^A \\ 0 & \|s_i - s_j\|^2 \leq \theta^A \end{cases} \quad (1.7)$$

The repulsive term is of the form

$$R_i(t) = \sum_{p_j \in r_i} \mu_j(t)(s_i - s - j)\delta_{ij}^R \quad (1.8)$$

where

$$\delta_{ij}^R = \begin{cases} \|s_i - s_j\|^{-2} & \|s_i - s_j\|^2 \leq \theta^R \\ 0 & \|s_i - s_j\|^2 > \theta^R \end{cases} \quad (1.9)$$

The two delta functions have different roles in the definition of the GS. The attractive δ^A corresponds to the inverse to the distance strength of attraction. To avoid singular values when two particles are close to zero distance we set a threshold θ^A which determines the region around the particles where the motion due to attraction forces disappear. The repulsive δ^R defines ij the maximum extension of the repulsive forces, which are short range forces. The threshold θ^R determines the region around the particles where the repulsive forces are active. A vertex particle of a GS-GC reaches zero velocity if and only if it is at distance below θ^A of a color particle and no repulsive particle is in θ^R range. A global state of the GS-GC is stationary if and only if all vertex particles are placed in the neighbourhood of some color particle without any repulsive particles located at the same color particle neighbourhood. If the graph's chromatic number M^* is smaller than or equal to the number of color particles $M^* \leq M$, there will be a non-empty set of stationary states of the GS-GC. Any stationary state of the GS-GC corresponds to a graph colouring. If the graph's chromatic number is greater than the number of color particles, there are no stationary states in the GS-GC.

These conditions mean that it is always possible to (given enough time) find the chromatic number for a graph. The algorithm to do this is outlined in the following sections.

1.4 Genetic Algorithm

1.5 Flower Pollination

The flower pollination algorithm[?] (FPA) was first described by Xin-She Yang¹. It is a metaheuristic global optimisation algorithm inspired by the pollination process of flowering plants. Although originally formulated to solve continuous optimisation problems, Meriem Bensouyad^{1.5} and Djamel Ed-dine Saidouni^{1.5} propose a discrete version of FPA for solving the graph coloring problem[?].

² Pollination can take two major forms: abiotic and biotic. About 90% of flowers belong to the class of biotic pollinating flowers, that is, their pollen is transferred by pollinators such as insects, birds, bats and other animals. The remaining 10% of flowers use abiotic pollination, which does not require any pollinators. Wind and diffusion in water help pollination of such flowering plants, of which grass is a good example. Pollinators, or sometimes called pollen vectors, can be very diverse. It is estimated that there are at least 200,000 varieties of pollinators.

A concept that comes up when talking about pollinators is *constancy*, which refers to the tendency of some pollinators to visit certain flower species while bypassing others. It has been conjectured that constancy has evolutionary advantages for both the flower species and the pollinator species. Pollination can be achieved by self-pollination or cross-pollination. Cross-pollination, or allogamy, means pollination can occur from pollen of a flower of a different plant, while self-pollination is the

¹Department of Engineering, University of Cambridge

²MISC Laboratory Constantine 2 University

fertilisation of one flower, such as peach flowers, from pollen of the same flower or different flowers of the same plant, which often occurs when there is no reliable pollinator available.

Biotic, cross-pollination may occur at long distances, and the pollinators such as bees, bats, birds and flies can fly a long distance, thus they can be considered as the global pollination. In addition, bees and birds may behave as Lévy flight behaviour, with jump or fly distance steps obey a Lévy distribution. Furthermore, flower constancy can be used as an incremental step using a similarity or difference of two flowers.

JESUE FUCK THESE PAPERS ARE BADLY WRITTEN

1.5.1 Flower Pollination Algorithm

The authors of the second paper do not deviate at this stage from the original, and both introduce the actual algorithm by arguing that the above characteristics can be idealised as follow:

1. Biotic and cross-pollination is considered as global pollination process with pollen-carrying pollinators performing Lévy flights.
2. Abiotic and self-pollination are considered to be local pollination.
3. Flower constancy can be considered as the reproduction probability is proportional to the similarity of two flowers involved.
4. Local pollination and global pollination is controlled by a switch probability $p \in [0, 1]$. Due to the physical proximity and other factors, such as wind, local pollination can have a significant fraction p in the overall pollination activities.

For simplicity's sake, the authors assume that each plant only has one flower and each flower only produces one pollen gamete³. In reality the number of flowers and pollen gametes per can vary widely between species and individual plants. This means that a solution, which we denote x_i , is equivalent to a flower and/or a pollen gamete⁴. The original author believes that extending the algorithm to multiple pollen gametes and multiple flowers (for multiobjective optimisation problems - I did not know that was a thing) should be easy.

The FPA is described by breaking it down into 2 key steps: global pollination and local pollination. For global pollination it is argued that the possibility of large travel distances for pollens and flower constancy give rise to the following step rule:

$$x_i^{t+1} = x_i^t + L(x_i^t - g_*) \quad (1.10)$$

where x_i^t is the pollen i or solution vector x_i at iteration t and g_* is the current best solution found amongst all solutions at the current iteration.

At this point we have to mention that the argument for global pollination being based on the current best solution, “the fittest flower”, is non-existent - it just appears in the middle of talking about travel distances and constancy without any justification.

³sperm cells

⁴There is no distinction between a flower and the pollen it produces, all that matters is how the pollen is used.

The parameter L is supposedly the “strength” of the pollination, which is a step size drawn from a Lévy distribution

$$L \frac{\lambda \gamma(\lambda) \sin(\pi \lambda / 2)}{\pi} \frac{1}{s^{1+\lambda}}, (s \gg s_0 > 0) \quad (1.11)$$

Here $\gamma(\lambda)$ is the standard gamma function, and this distribution is valid for large steps $s > 0$. Yang reports using $\lambda = 1.5$ for the simulations in the original paper, the second paper’s authors (for GCP) do not report what value they used⁵.

The local pollination, which also includes an allowance for constancy apparently, can be represented as

$$x_i^{t+1} = x_i^t + \epsilon(x_j^t - x_k^t), \quad (1.12)$$

where x_j^t and x_k^t are pollens from the different flowers of the same plant species and ϵ is drawn from a uniform distribution in $[0, 1]$. (WE MAY HAVE FUCKED THIS UP? WTF IS THE SAME PLANT SPECIES IN THIS ALGORITHM?)

Most flower pollination activities can occur at both local and global scale. In practice, adjacent flower patches or flowers in the not-so-far-away neighbourhood are more likely to be pollinated by local flower pollens than those far away. For this, we use a switch probability p to switch between common global pollination and intensive local pollination. (I do not know what “common” and “intensive” mean, they are never explained.) In the original paper (Yang) the author describes using an initial $p = 0.5$ and then performing a parametric study which found that $p = 0.8$ worked best for most applications. Again, the authors of the second paper did not discuss actual values.

⁵we will talk about why this is unhelpful later

Chapter 2

Implementations & Original Contributions

2.1 Introduction

This section outlines specific implementations of some of the algorithms described above. Where possible pseudocode is provided. In most cases it is *not* the complete algorithm and further analysis of the original paper will be required to reconstruct the algorithm. Results for comparison of the algorithms is given in the Results section.

2.2 Random Brute Buckets

The Random Brute Buckets (RBB) algorithm presented here is a class of Successive Augmentation Technique. It is at its core an implementation of an insertion sort algorithm which samples vertices from the graph in a random order. This algorithm is guaranteed to find a valid colouring for graphs that are directed and undirected. It is however not bounded, and there is no guarantee that the algorithm will converge upon the best colouring.

2.2.1 Implementation and Problems

The RBB algorithm is presented in Algorithm 1 below.

Algorithm 1 Random Brute Buckets

```
1: procedure SOLVE(Graph  $g$ , long  $iterationLimit$ ) ▷  $g$  is predefined
2:    $currentIteration \leftarrow 0$ 
3:    $currentBestColouring \leftarrow \infty$ 
4:   Create empty list of buckets  $bList$ 
5:   while  $currentIteration < iterationLimit$  do
6:     Populate vertex set  $V$ 
7:     while  $V \neq \emptyset$  &&  $|bList| < currentBestColouring$  do
8:        $v \leftarrow$  random vertex from  $V$  ▷ Remove  $v$  from  $V$ 
9:        $vertexPlaced \leftarrow FALSE$ 
10:      for all  $bucket \in bList$  do
11:        if  $bucket$  contains no conflicts with  $v$  then
12:          add  $v$  to  $bucket$ 
13:           $vertexPlaced \leftarrow TRUE$ 
14:          break for
15:        else continue
16:      end if
17:    end for
18:    if  $!vertexPlaced$  then ▷ create a place for the vertex
19:      create new bucket  $b_0$ 
20:      add  $v$  to  $b_0$ 
21:      add  $b_0$  to  $bList$ 
22:    end if
23:  end while
24:  if  $|b| < currentBestColouring$  then
25:     $currentBestColouring = |bList|$ 
26:  end if
27:   $currentIteration ++$ 
28: end while
29: return  $currentBestColouring$ 
30: end procedure
```

This algorithm always improves upon the best colouring (which is initially set to ∞ on line 3) in the first iteration, and upon subsequent iterations, if a better solution is stumbled upon randomly, that solution is set as the current best solution. The algorithm returns the best solution found after a set number of iterations set at runtime.

2.2.2 Workarounds

This algorithm is very fast at finding initial solutions. This algorithm does not try to improve upon solutions that it finds, so local minima offer no resistance to finding the solution. This algorithm is not guaranteed to find the best solution. Indeed even after infinite attempts, there is no guarantee that the best solution will be found. That being said, the algorithm is generally able to find solutions close to the optimal colouring for most graphs extremely quickly and with a minimum use of resources (compared to other algorithms presented in this paper).

2.3 Gravitational Swarm Intelligence

The Gravitational Swarm Intelligence Algorithm presented in this paper is a modified version of the algorithm presented in the paper "Gravitational Swarm for Graph Coloring" by Israel Carlos Rebollo Ruiz [1]. Swarm Intelligence is a model where the emergent collective behavior is the outcome of a process of self-organization, where the agents evolve autonomously following a set of internal rules for their motion, interaction with the environment and the other agents. In this algorithm, the agents are subjected to an environment subject to a version of Newtonian Gravity. Centres of attraction (wells) are placed in the environment and the agents move based upon their attraction to these wells. Each well represents a distinct grouping or colouring of the underlying graph. Agents experience a repulsive force if they try and enter a well already containing agents that prohibit a valid colouring. If an agent enters a well that is not prohibited, then it stops moving and takes on that well's colour. The algorithm ends once a set iteration limit is reached, or all agents settle into the wells, and hence a valid colouring is found.

2.3.1 Implementation and Problems

The action of each agent on every iteration is presented below:




Figure 2.1: Agent logic for Gravitational Swarm

Each agent moves in the toric world until all agents fall into the "Stand By" State. Then

once the last agent triggers the "All Finished" State, the current colouring is accepted as valid. This algorithm presented some issues when implemented for testing. The comfort statistic that each agent tracks can grow unboundedly if the agents repeatedly fall into wells that they are subsequently repulsed from. This causes all captured agents to grow in comfort, making the local minima harder to improve with every iteration and the chance of jumping out and finding a valid solution fall dramatically.

The distance functions that were implemented were found to contain an error, resulting in a toroidal universe but gravitational attraction only in the plane. This meant that all results presented here use the non-toroidal distance. This does not present an issue, as the algorithm is still shown to converge to the solution as long as the distance functions work partially.

The figure below shows the location of wells within the universe, with distance functions permitting a toroidal universe. The colouring is a result of starting an agent at every possible start location and logging which well captures it.




Figure 2.2: Gravity mapping with gravity well locations overlaid

2.3.2 Workarounds

This algorithm was relatively easy to implement, however much tweaking was required to ensure that the algorithm converges in a reasonable time and iteration count. The strength of the "gravity" had to be varied, and different radii of capture were compared, as these factors were the factors that caused the algorithm's runtime to be most affected.

This algorithm was implemented from the psuedocode provided in [1], however the sample code omitted some details of the algorithm and instead presented these details the "Results and Observations" section of the paper. This lead to some aspects of the algorithm not functioning as intended.

Interestingly, once the algorithm was tweaked to function correctly, the algorithm reflected the same observations that the original author had made.

2.3.3 Additions

Substantial time and effort has been committed to the algorithm, with many aspects and details changing and growing over time to better solve the problems presented. The algorithm was trialled with the option for gravity wells to be able to take some "mass" from captured agents and have their own velocity in the universe. This addition did not increase the convergence rate of the algorithm but it did expose some flaws with the distance functions used to allow agents to track the wells.

Agent velocity and direction calculations have been changed and tweaked numerous times in the attempt to have agents move in the toroidal universe and react correctly to having multiple sources of strong attraction. To this end, a function was created that mapped agent start locations to final end points. Any "inconsistencies" are easily visually identified. An example of inconsistent behaviour is shown here:




Figure 2.3: Toroidal Gravity Map with inconsistent behaviour - top left, top right exhibit "water-drop" pattern

The large amounts of variability around disputed boundaries is a result of the random noise added to the velocity on every iteration. It was found that if no noise was added, when there was very low gravity values (an agent was a long distance from *any* well), the agent would stall, taking many iterations to gain speed and be captured. The random noise helps to speed up the process, and can reduce the amount of iterations in a stall by "randomly walking" around the region until the gravity values become larger than the random noise.

2.4 Genetic Algorithm

The Genetic Algorithm presented here is a variant of the algorithm described by [2]. The algorithm was implemented as described, then tweaked to improve efficiency and to experiment with parameter tweaking and the effect of changing certain flows within the algorithm. These changes are discussed further in the Additions (2.4.3) section.

2.4.1 Implementation and Problems

The algorithm is implemented as follows:

Algorithm 2 Genetic Algorithm with Wisdom of Crowds

```
1: procedure SOLVE(Graph  $g$ ,  $iterationLimit$ ,  $numChromosomes$ )  $\triangleright g$  is predefined
2:    $currentAttempt \leftarrow 0$ 
3:    $currentBestColouring \leftarrow \Delta(g) + 1$ 
4:    $aggregateChromosome \leftarrow$  chromosome of randomly assigned colours.
5:   while  $currentIteration < iterationLimit$  do
6:      $population \leftarrow$  set of chromosomes with randomly assigned colours. (up to  $numColours$ )
7:     if solvePop() then
8:        $aggregateChromosome \leftarrow$  chromosome of randomly assigned colours.
9:        $currentAttempt \leftarrow 0$ 
10:    else
11:       $currentAttempt \leftarrow currentAttempt + 1$ 
12:    end if
13:  end while
14:  return  $currentBestColouring$ 
15: end procedure
```

Algorithm 3 Genetic Algorithm with Wisdom of Crowds - Tick Generation

```
1: procedure SOLVEPOP(Graph  $g$ ,  $iterationLimit$ ,  $numChromosomes$ )  $\triangleright g$  is predefined
2:    $currentIteration \leftarrow 0$ 
3:   while  $currentIteration < iterationLimit$  and best solution has cost  $> 0$  do
4:      $currentIteration \leftarrow currentIteration + 1$ 
5:     if best chromosome has cost  $\geq altMethodThreshold$  then
6:        $parents \leftarrow getParentsA()$ 
7:     else
8:        $parents \leftarrow getParentsB()$ 
9:     end if
10:     $child \leftarrow crossOver(parents)$ 
11:    if  $rand < mutChance$  then
12:      if best chromosome has cost  $\geq altMethodThreshold$  then
13:         $child \leftarrow mutateA()$ 
14:      else
15:         $child \leftarrow mutateB()$ 
16:      end if
17:    end if
18:    add  $child$  to  $population$ 
19:    remove bottom performing half of population
20:    repopulate up to  $numChromosomes$ 
21:  end while
22:  if  $currentIteration \geq iterationLimit$  then
23:    perform  $wisdomOfCrowds()$   $\triangleright$  generate  $aggregateChromosome$  by voting
24:    add  $aggregateChromosome$  to population
25:  end if
26:  if best solution has cost 0 then
27:     $currentBestSolution \leftarrow currentBestSolution - 1$ 
28:    return true
29:  else
30:    return false
31:  end if
32: end procedure
```

Algorithm 4 Genetic Algorithm with Wisdom of Crowds - Parent Selection

```
1: procedure GETPARENTSA
2:    $tempParents \leftarrow$  choose two random chromosomes from population.
3:    $parent1 \leftarrow$  fitter of  $tempParents$ 
4:    $tempParents \leftarrow$  choose two random chromosomes from population.
5:    $parent2 \leftarrow$  fitter of  $tempParents$ 
6:   return  $parent1, parent2$ 
7: end procedure
8:
1: procedure GETPARENTSB
2:   return top performing chromosome, top performing chromosome
3: end procedure
```

Algorithm 5 Genetic Algorithm with Wisdom of Crowds - Crossover

```
1: procedure CROSSOVER
2:    $child \leftarrow$  colours up to and including a random point from  $parent1$ , followed by the colours
   from  $parent2$  from that point on in the chromosome.
3:   return child
4: end procedure
```

Algorithm 6 Genetic Algorithm with Wisdom of Crowds - Child Mutation

```
1: procedure MUTATEA
2:   for all  $vertex$  in chromosome do
3:     if  $vertex$  has a conflict then
4:        $adjacentColours \leftarrow$  all adjacent colours to  $vertex$ 
5:        $validColours \leftarrow allColours - adjacentColours$ 
6:        $newColour \leftarrow$  random colour from  $validColours$ 
7:       set chromosome colour at  $vertex$  to be  $newColour$ 
8:     end if
9:   end for
10: end procedure
11:
1: procedure MUTATEB
2:   for all  $vertex$  in chromosome do
3:     if  $vertex$  has a conflict then
4:        $newColour \leftarrow$  random colour from allColours
5:       set chromosome colour at  $vertex$  to be  $newColour$ 
6:     end if
7:   end for
8: end procedure
```

Algorithm 7 Genetic Algorithm with Wisdom of Crowds - Wisdom Of Artificial Crowds

```
1: procedure WISDOMOFCROWDS
2:    $expertChromosomes \leftarrow$  best half of final population
3:    $aggregateChromosome \leftarrow$  best performing chromosome
4:   for all  $vertex$  do
5:     if  $vertex$  is part of a bad edge then
6:        $newColour \leftarrow$  the most used colour for  $vertex$  among  $expertChromosomes$ 
7:       set colour at  $vertex$  of  $aggregateChromosome$  to be  $newColour$ 
8:     end if
9:   end for
10: end procedure
```

2.4.2 Workarounds

The algorithm provides some parameters to tailor the solving method to be able to deal with local optima in several ways. Firstly there are 2 different parent selection methods that are used depending upon how close the best chromosome is to a valid colouring. Secondly there are two different mutation

methods that are used depending upon how close the best chromosome is to a valid colouring. It was found that while these methods allowed for some level of control over the algorithm, it was by no means as configurable as the other algorithms that were implemented. To remedy this and to produce our own variant of the Genetic Algorithm for solving the GCP, alterations were made to the flow to allow more control over the algorithm.

2.4.3 Additions

The first deviation to the documented algorithm was to the parent selection criterion. It was found that having the best solution be the only parent that is mutated when it is sufficiently close to a valid colouring slowed the convergence rate of the other chromosomes to zero. To remedy this, we implemented a switch such that a certain percent of the time, the parents chosen were the best chromosome and itself. But a small percentage of the time, the parents would be the best chromosome and a random chromosome from the population. This prevented the stagnation of the problem, and allowed for other changes to be made to improve the algorithm. The second part of the algorithm identified for change was the crossover method. This method originally selected the first portion of *parent1* and mixed it with the complement of that section of *parent2*. This process was changed to a simple 50/50, such that the chance of the first section coming from *parent1* was equally weighted to the first section coming from *parent2*. This prevented possible bad colourings from being propagated simply because they occurred early in the chromosome. The final change to the algorithm was in response to results collected from the algorithm. It was found the wisdom of crowds voting mechanism was finding the solution far more successfully than the Genetic Algorithm itself. This prompted a change to the algorithm whereby a vote could not take place before certain conditions were met in the agreement of the chromosomes. This led to a much higher runtime for the algorithm, but once a vote could commence, the algorithm was much more likely to either improve the colouring or produce a chromosome with many fewer conflicts. This method was settled upon after several other methods were trialed and discarded. As the algorithm searches for a tighter colouring, it takes many more generations for the chromosomes to converge. This method was settled upon since it allowed the algorithm to be applied to any graph with a minimum of tweaking, yet a valid vote will always take place, since the algorithm will always reach a point of agreement, even if the solution is not valid for the given graph. Setting the level of agreement that each chromosome must exhibit allows the algorithm to be sped up or slowed down. If the limit is relaxed, each trial takes less time to complete, yet the wisdom of crowds vote has a lower chance of improving the solution. Tightening the bounds has the effect of dramatically increasing the runtime of the algorithm but increases the chance that the wisdom of crowds vote will yield a better solution. Setting the bound too tight however will result in the algorithm running for an infinite amount of time (if the level of agreement specified is not possible). The last change that was made to this algorithm was to create a variable parameter for how much of the population was discarded and repopulated at random on each generation. This parameter is responsible for setting how much of the solution space is tried. Each iteration, the random nature of the generation process has the possibility to generate a better solution to the problem. However, the larger the proportion of the population that is discarded upon each iteration, the lower probability that the wisdom of crowds vote that occurs at the end of each trial will improve upon the solution. For this reason, the proportion of the population that is allowed to participate in the vote is set to always be less than (or equal to) the proportion that is kept in each iteration, so that random noise is not added to the voted data when a vote takes place.

2.5 Flower Pollination

2.5.1 Implementation and Problems

2.5.2 Workarounds

2.5.3 Additions

Chapter 3

Results & Discussion

3.1 Algorithms

Graph	χ	$\Delta + 1$	Solver	Limits on Runtime	k	time to k	% success
queen5_5.col	5	del+1	RBB	100000	5	time to k	percent
queen5_5.col	5	del+1	GS	1000000	5	time to k	percent
queen5_5.col	5	del+1	FPA	10000	7	time to k	percent
queen5_5.col	5	del+1	GA	20000	7	time to k	percent
myciel5.col	6	del+1	RBB	100000	3	time to k	percent
myciel5.col	6	del+1	GS	1000000	2	time to k	percent
myciel5.col	6	del+1	FPA	10000	10	time to k	percent
myciel5.col	6	del+1	GA	20000	8	time to k	percent
queen7_7.col	7	del+1	RBB	100000	7	time to k	percent
queen7_7.col	7	del+1	GS	1000000	9	time to k	percent
queen7_7.col	7	del+1	FPA	10000	19	time to k	percent
queen7_7.col	7	del+1	GA	20000	-	-	percent
mulsol.i.3.col	31	del+1	RBB	100000	31	2844	percent
mulsol.i.3.col	31	del+1	GS	1000000	31	time to k	percent
mulsol.i.3.col	31	del+1	FPA	10000	-	-	percent
mulsol.i.3.col	31	del+1	GA	20000	-	-	percent

Table 3.1: Results showing the time taken for the algorithm to find a valid colouring (k colours), optimal or not

Graph	χ	$\Delta + 1$	Solver	Limits on Runtime	min k reached	time to k	total runtime
queen5.5.col	5	del+1	RBB	100000	5	time to k	395.7
queen5.5.col	5	del+1	GS	1000000	5	time to k	1382389
queen5.5.col	5	del+1	FPA	10000	7	time to k	2136729
queen5.5.col	5	del+1	GA	20000	7	time to k	7211935
myciel5.col	6	del+1	RBB	100000	3	time to k	470.3
myciel5.col	6	del+1	GS	1000000	2	time to k	1.63×10^7
myciel5.col	6	del+1	FPA	10000	10	time to k	4899871.0
myciel5.col	6	del+1	GA	20000	8	time to k	2008825
queen7.7.col	7	del+1	RBB	100000	7	time to k	780.0
queen7.7.col	7	del+1	GS	1000000	9	time to k	1.23×10^7
queen7.7.col	7	del+1	FPA	10000	19	time to k	13875131
queen7.7.col	7	del+1	GA	20000	-	-	-
mulsol.i.3.col	31	del+1	RBB	100000	31	2844	4952
mulsol.i.3.col	31	del+1	GS	1000000	31	time to k	21604849
mulsol.i.3.col	31	del+1	FPA	10000	-	-	-
mulsol.i.3.col	31	del+1	GA	20000	-	-	-

Table 3.2: Results of algorithms attempts at finding a better solution (and failing)

Bibliography

- [1] Israel Carlos Rebollo Ruiz. *Gravitational Swarm for Graph Coloring*. <http://www.chu.eus/ccwintco/uploads/3/3a/Tesis-Israel-final.pdf>, 2012.
- [2] Musa M. Hindi and Roman V. Yampolskiy. *Genetic Algorithm Applied to the Graph Coloring Problem*. http://ceur-ws.org/Vol-841/submission_10.pdf, 2015.